

| Cloud Native Aarhus



# MAKE MONEY MATTER.

By Kasper Nissen (@phennex), DevOps Engineer @thelunarway



| Cloud Native Aarhus




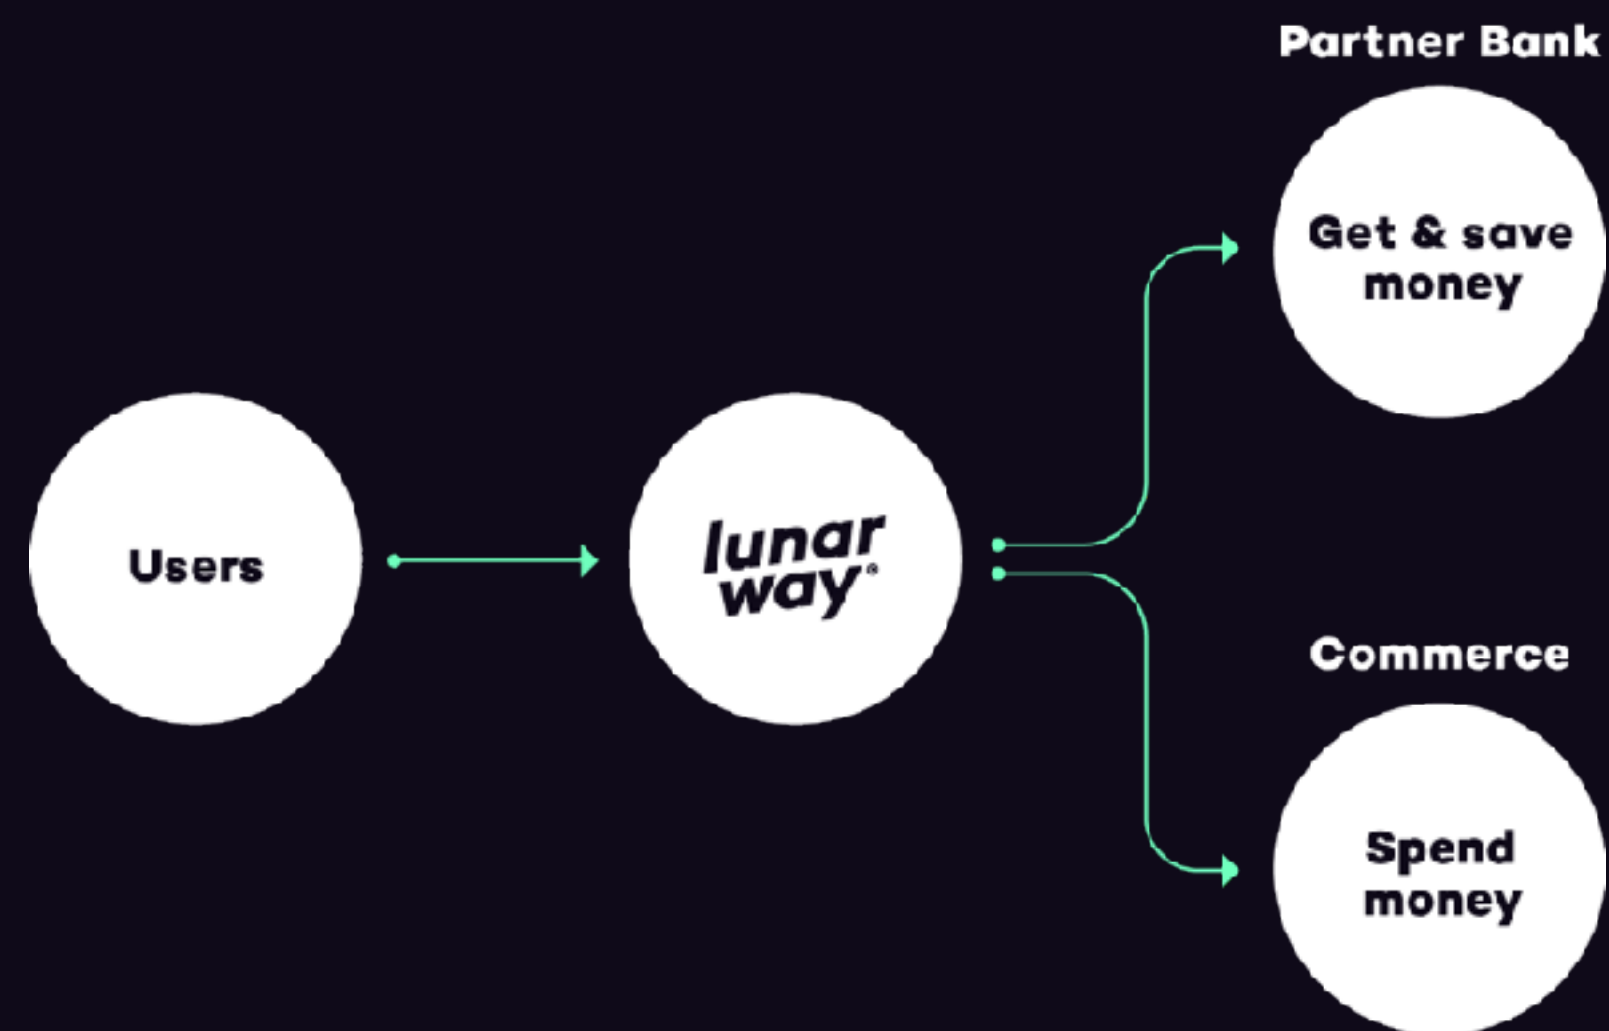
# MAKE MONEY MATTER.

By Kasper Nissen (@phennex), DevOps Engineer @thelunarway



# ***Lunar Way***

- The Partner model
  - Leverage the partner banks infrastructure
  - All money is in the partner bank
- Currently only in Denmark, will move to the nordics in the near future





# Kasper Nissen

DevOps & Infrastructure Engineer @thelunarway

## Experience

DevOps & Infrastructure Engineer @ LEGO (CITMABIS) (oursourced by IT Minds) for 5 months

Senior/Software Engineer @ IT Minds (~4 years part time)

Master thesis: KubeCloud - A Small-Scale Tangible Cloud Computing Environment.

Interview with Software Engineering Daily: [bit.ly/2paZ5lg](https://bit.ly/2paZ5lg)

Blogging about Cloud Native Tech @ [www.kubecloud.io](http://www.kubecloud.io)

M. Eng. Computer Technology from Aarhus University - Department of Engineering.

B. Eng. Information and Communication Technology from Aarhus University - School of Engineering





# ***What do we have running?***

**19**  
services

**215**  
containers  
in prod

**3**  
kubernetes  
clusters

**1**  
rails  
monolith

**3**  
node rabbitmq  
cluster

**2**  
100 GB  
postgresql  
DB's


**3**  
AWS Accounts

**13**  
infrastructure  
services

***Where are we running?***



# ***Service overview***



# Service overview *(with infrastructure)*






# ***Infrastructure***

***AWS***



***RabbitMQ Cluster***



***Kubernetes Cluster***



***Elasticsearch  
Cluster***




***PostgreSQL  
database***


# ***Where are we heading?***

## ***Horses vs Unicorns***


- Gene Kim




# ***Cloud Native Utopia***



***Microservice oriented***




***Container packaged***




***Dynamically scheduled***





# ***MICROSERVICES***



# ***Why ?***

*(our Microservice vision)*

## **1.**

- ***Development***

- Freedom and autonomy
- Best tool for the job
- Speed

## **2.**

- ***Architecture***

- Fault tolerance
- Flexibility
- Coherence, decoupling, encapsulation

## **3.**

- ***Deployment***

- Independence
- Scalability
- Speed
- Resource utilisation

# ***How are we building our services?***

## ***Asynchronous first***




Decoupling in time and space allows for autonomy

## ***Shared dependencies***

Common packages, such as logging, monitoring, communication




## ***Each service has it's own repository***

Containing source code, deployment spec, pipeline, etc.







# ***Killing our Rails monolith***



*... slowly strangling Rails*

# ***Languages***



# ***Challenges - so far***

- ***Development***

- Tooling - e.g. Swagger for events?
- Local env
- Shared code

- ***Deployment***

- Automation
- Versioning

- ***Architecture***


- Distributed monolith
- Transparency
- Complexity
- Cross cutting concerns

- ***Operation***

- Monitoring
- Tracing events



# ***CONTAINERS***



# ***What?***

***DEVELOPERS***

***Container***



***App***

***Runtime***

***OS***

***HOST OS***

***OPERATIONS***



# Why?

## Isolation


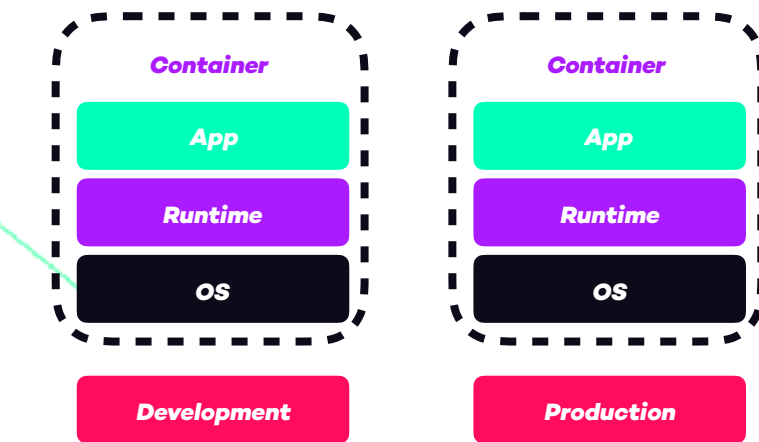

Services are isolated and contained in their environment

## Consistency in portability

The container will run in the same way in local env as in prod

## Versioning

Versioning a container is easy, rolling back and forth becomes easy








The image features a dark blue background with the text "DYNAMICALLY SCHEDULED" in a bold, italicized, white sans-serif font. The text is centered horizontally. On the left side, there is a vertical stack of six teal rectangular blocks, each containing ten vertical white lines. On the right side, there is a larger, more complex arrangement of similar teal blocks, some stacked vertically and others horizontally, creating a stepped effect. The overall aesthetic is modern and technical.

***DYNAMICALLY  
SCHEDULED***

# ***What?***



# ***Why?***

## ***Scheduling***

The scheduler will schedule your service on a node

## ***Resource optimization***

Scheduling allows for better packaging of services in hosts

## ***Resiliency***

If a node dies, the scheduler will reschedule on another node

## ***Scalability***

Scaling a dynamically cluster is easy, just add more nodes

**SO,**  
**HOW DO WE WORK**  
**TOWARDS BECOMING**  
**A UNICORN?**



# ***Currently used CNCF projects***




OPENTRACING



fluentd

GRPC



linkerd




containerd



rkt

***Container Orchestration with  
Kubernetes***



# ***Why Kubernetes?***

## ***Community***

48k+ commits, 22+ GitHub stars, 1.1k contributors

## ***Crossplatform***

Multiple arcs, multiple cloud providers

## ***Resource optimization***

Packing nodes to utilize available resources

## ***Tooling***

A lot of great tools

## ***Scaling***

Great integration with auto-scaling, both on node- and container-level

## ***High availability***

Automatic failover, redundancy

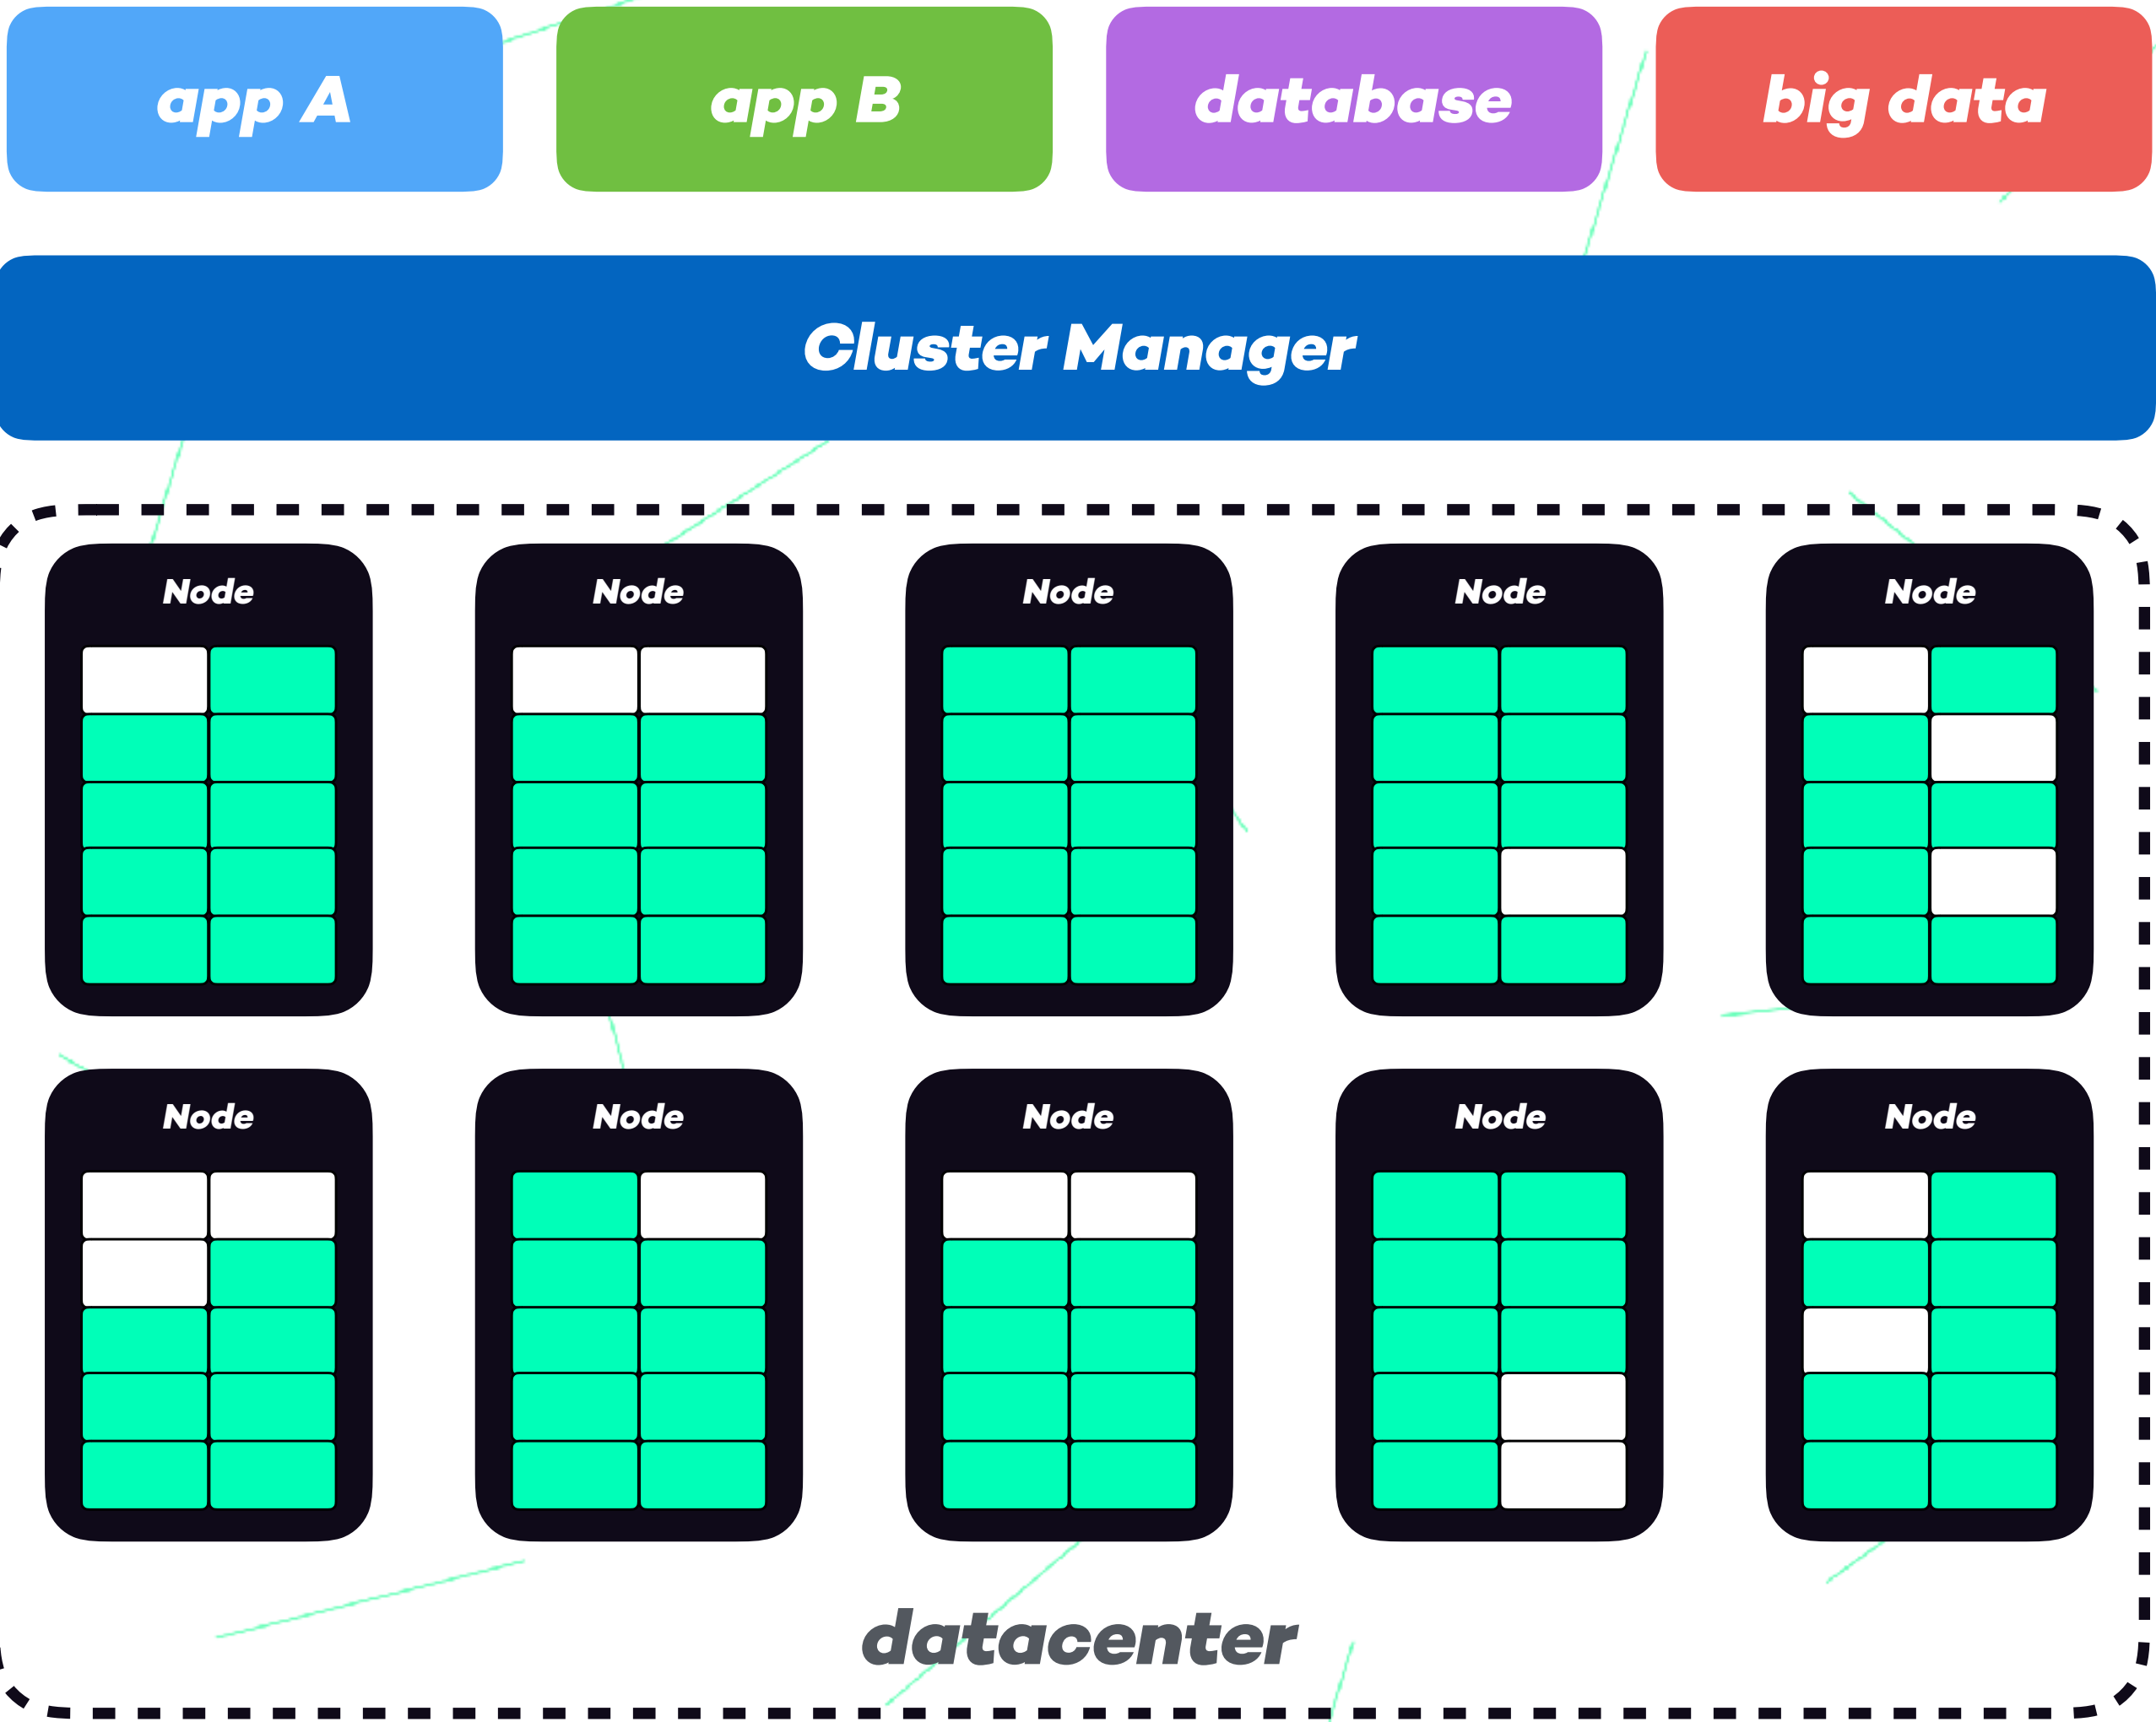
## ***Pure Awesomeness!***

It's just awesome!

**1.5 release:**  
Estimated 400 years of work hours

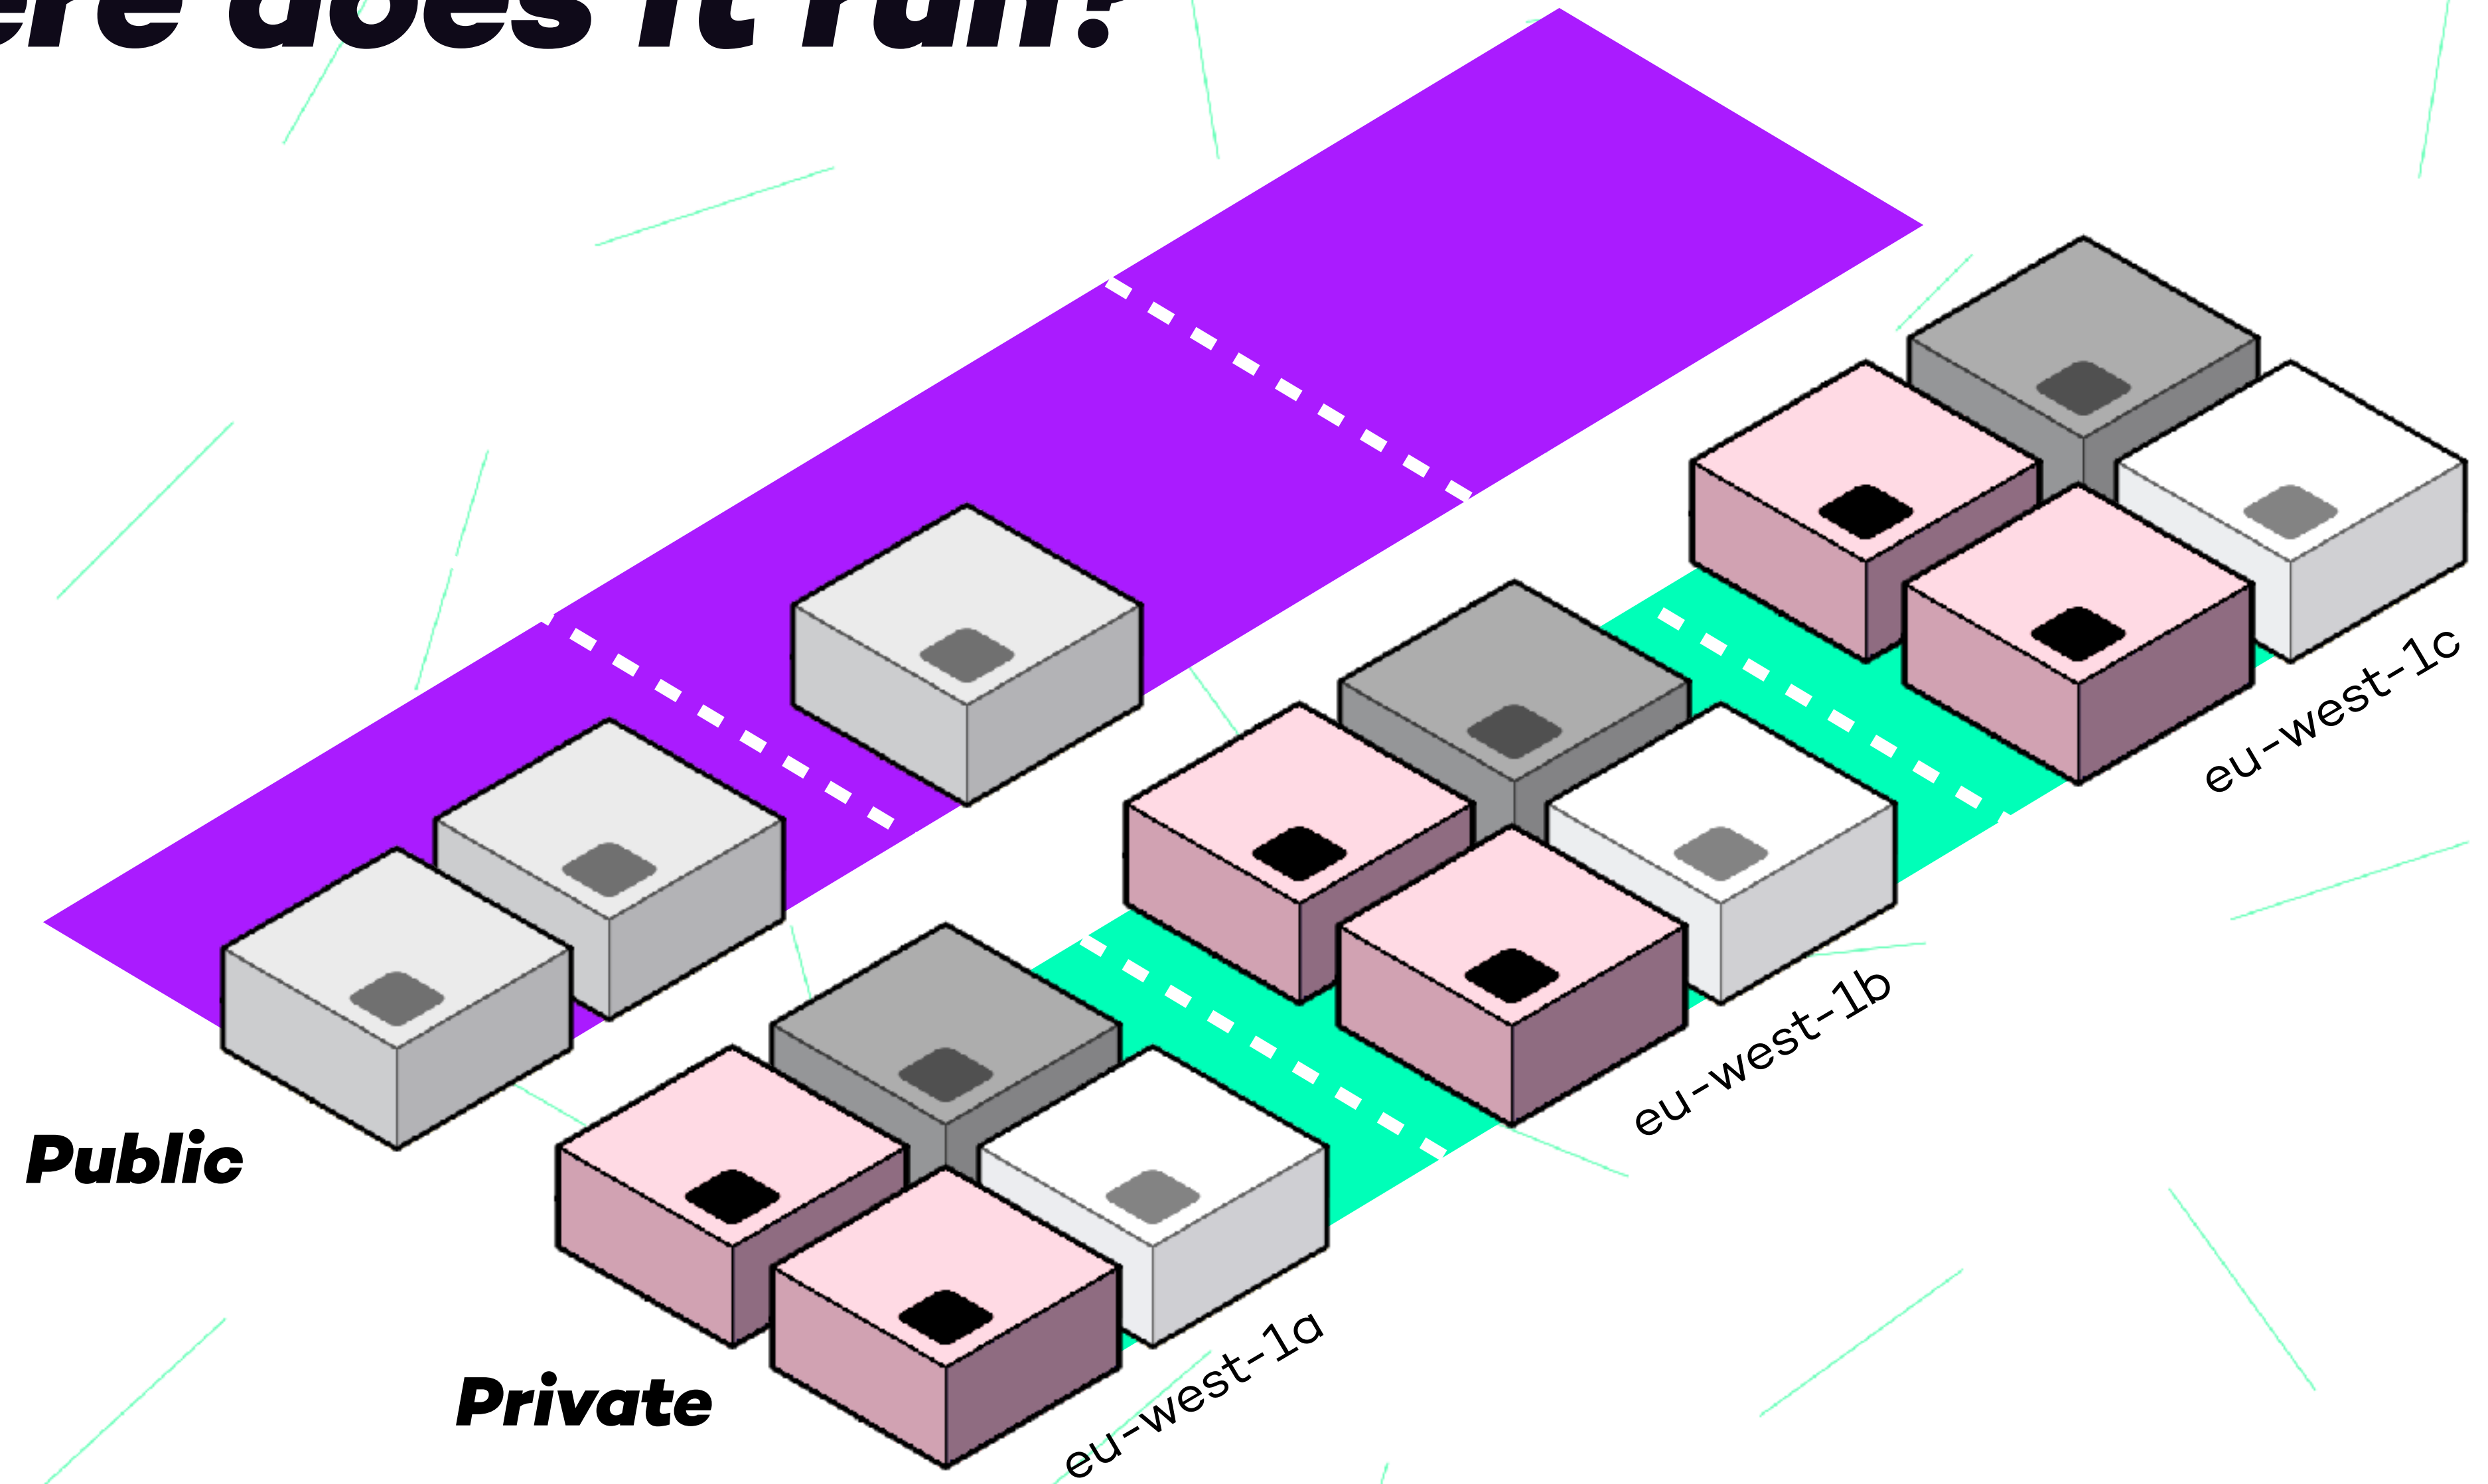


# What does it do?

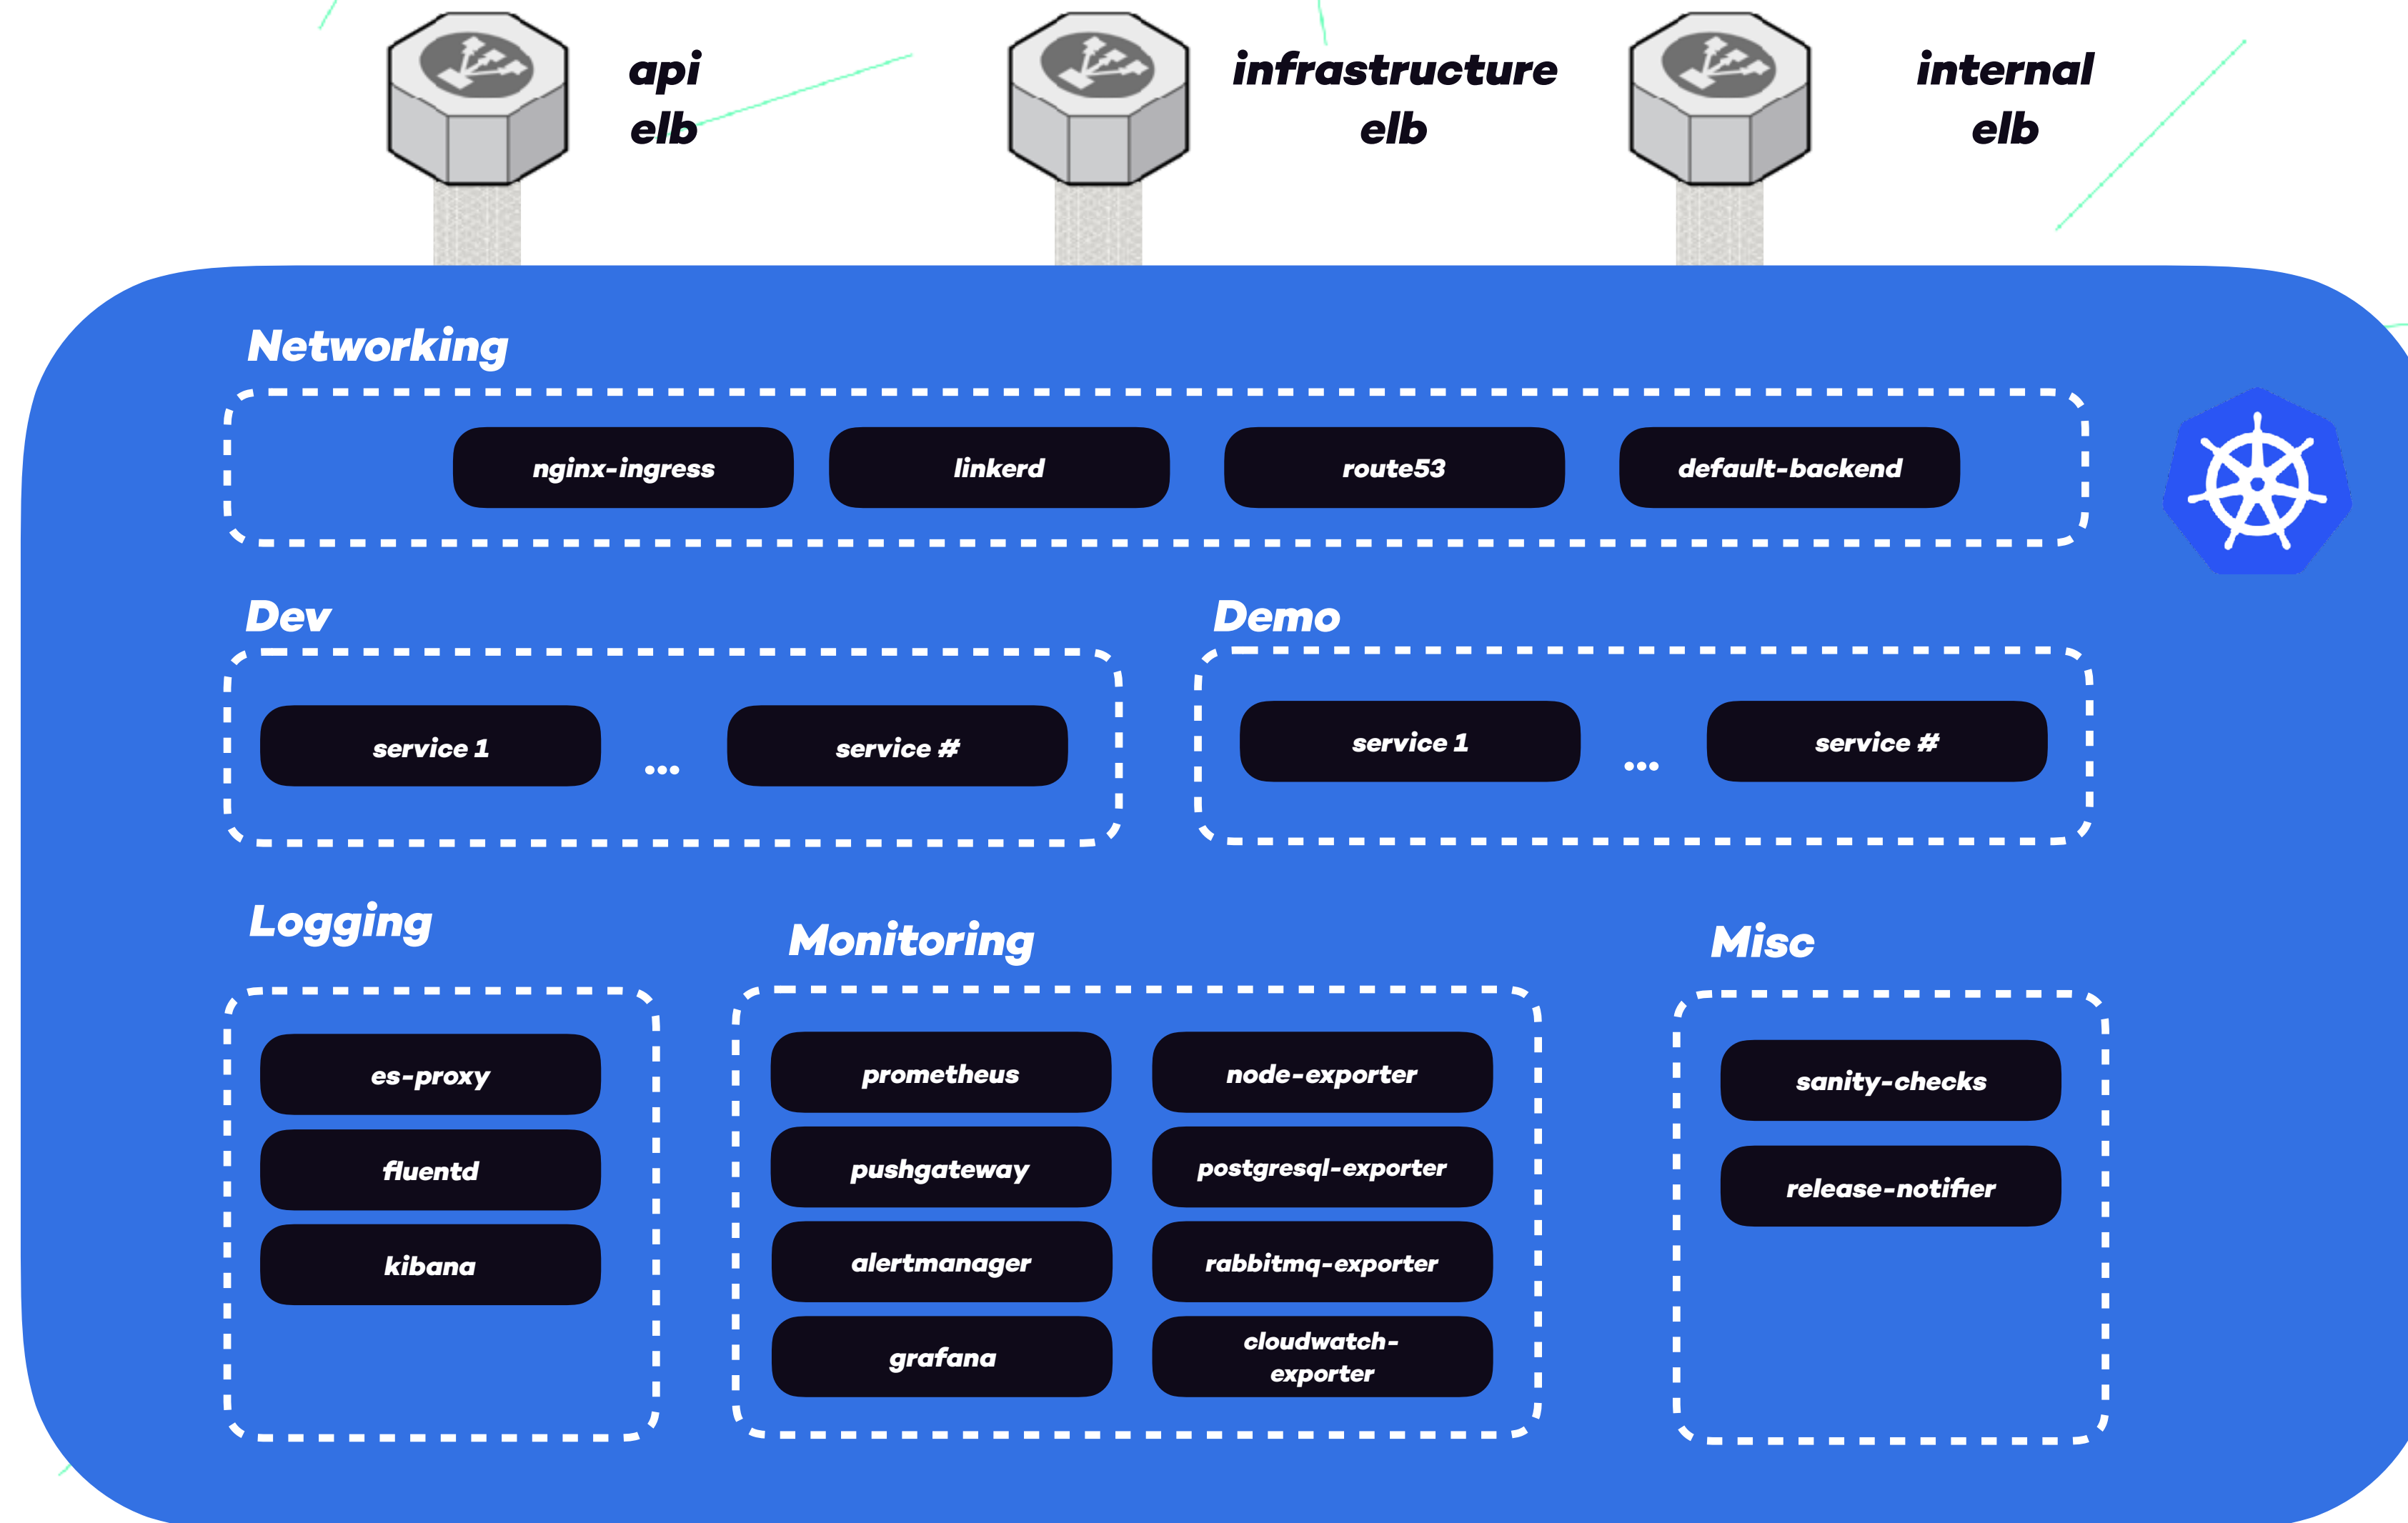




# ***Where does it run?***



# Services running in kubernetes



# ***What do we think of it?***

## ***Freedom***

Squads can deploy and more or less implement how they see fit

## ***Autonomous services***

Squads can work independent of other squads

## ***Continuous Delivery***

Kubernetes allows us to deploy multiple times a day. It's easy to rollback in case something went wrong

## ***Flexibility***

We run many different type of workloads in the cluster. Gives us mobility to become cloud agnostic

## ***Scalable infrastructure***

Scaling the infrastructure is easy, both on node and container level

## ***High availability***

Kubernetes takes care of container failures, AWS Auto Scaling groups takes care of node failures

## ***Easy maintenance***

We are using Kubernetes Operations to help us spin up our clusters, and maintain them.



*Log collection with*  
***Fluentd***



# ***Why fluentd?***

## ***Simple and Easy***

Provides a simple interface for specifying input and output. Works great with Kubernetes and containers.

## ***Community***

Big community around fluentd, validates our choice.

## ***Small memory footprint***

Do not require a lot of resources in the cluster

## ***Proven reliability and performance***

It's a fairly battle tested project

# What does it do?

1.

```
<source>
  @type tail
  path /var/log/containers/*.log
  pos_file /var/log/containers/es-containers.log.pos
  tag kubernetes.application.*
  format json
  time_key event_time
  time_format %Y-%m-%dT%H:%M:%S.%3NZ
</source>
```


2.

```
<filter kubernetes.application.**>
  @type kubernetes_metadata
  merge_json_log true
  preserve_json_log false
</filter>
```

3.

```
<match kubernetes.application.**>
  type elasticsearch
  host es-proxy
  port 9200
  include_tag_key true
  logstash_format true
  logstash_prefix application-
  reload_on_failure true
</match>
```

# ***Logging setup***



***AWS Elasticsearch Cluster***

# ***What do we think of it?***

## ***Very easy to use***

Set up is easy!

## ***Works great with Kubernetes***

Awesome plugin for adding Kubernetes metadata - making it easy to identify pods etc.

## ***Run as a daemonset***

Easy to run in every node of the cluster as a daemonset



***Monitoring with***  
***Prometheus***



# What?

*short-lived jobs*

*long-lived jobs*


**Pushgateway**

scrapes

**Prometheus**

**Grafana**

**Alertmanager**



# ***What does it do?***

## ***Multi-dimensional data model***

Time series identified by metric name and key/value pairs

## ***Flexible query language***

Comes with a builtin query language for all kinds of operations, sums, averages, increase, rate, etc.

## ***Easy and simple***

Easy to setup, and works great with Kubernetes service discovery

## ***Alerting and great integration with Grafana***

Prometheus has a builtin alerting system, and Grafana provides easy integration for making metrics visible in dashboards

## ***Pull-based approach***

Prometheus scrapes its targets at a regular interval

# ***What metrics are we collecting?***

## ***Kubernetes specific metrics***

Pods running, health of Kubernetes system components, etc.

## ***RabbitMQ***

Activity in queues, unacknowledged messages

## ***Nodes***

CPU, Memory

## ***Traffic***

Incoming traffic, upstream latency in cluster, etc.

## ***Containers***

CPU, Memory

## ***Application Specific metrics***

Relevant metrics, instrumented by the services owners



***DEMO***



# ***What do we think of it?***

## ***Provides great insights***

Provides valuable insights in the state of the cluster

## ***Makes is easy to developers to instrument their services***

We provide a simple package for instrumentation, making squads able to do their own monitoring. ***YOU BUILT IT, YOU RUN IT!***

## ***Grafana integrations is sweet!***

Grafana and Prometheus works well together, making Grafana the interface for building dashboards and alerts

## ***Kubernetes <3 Prometheus***

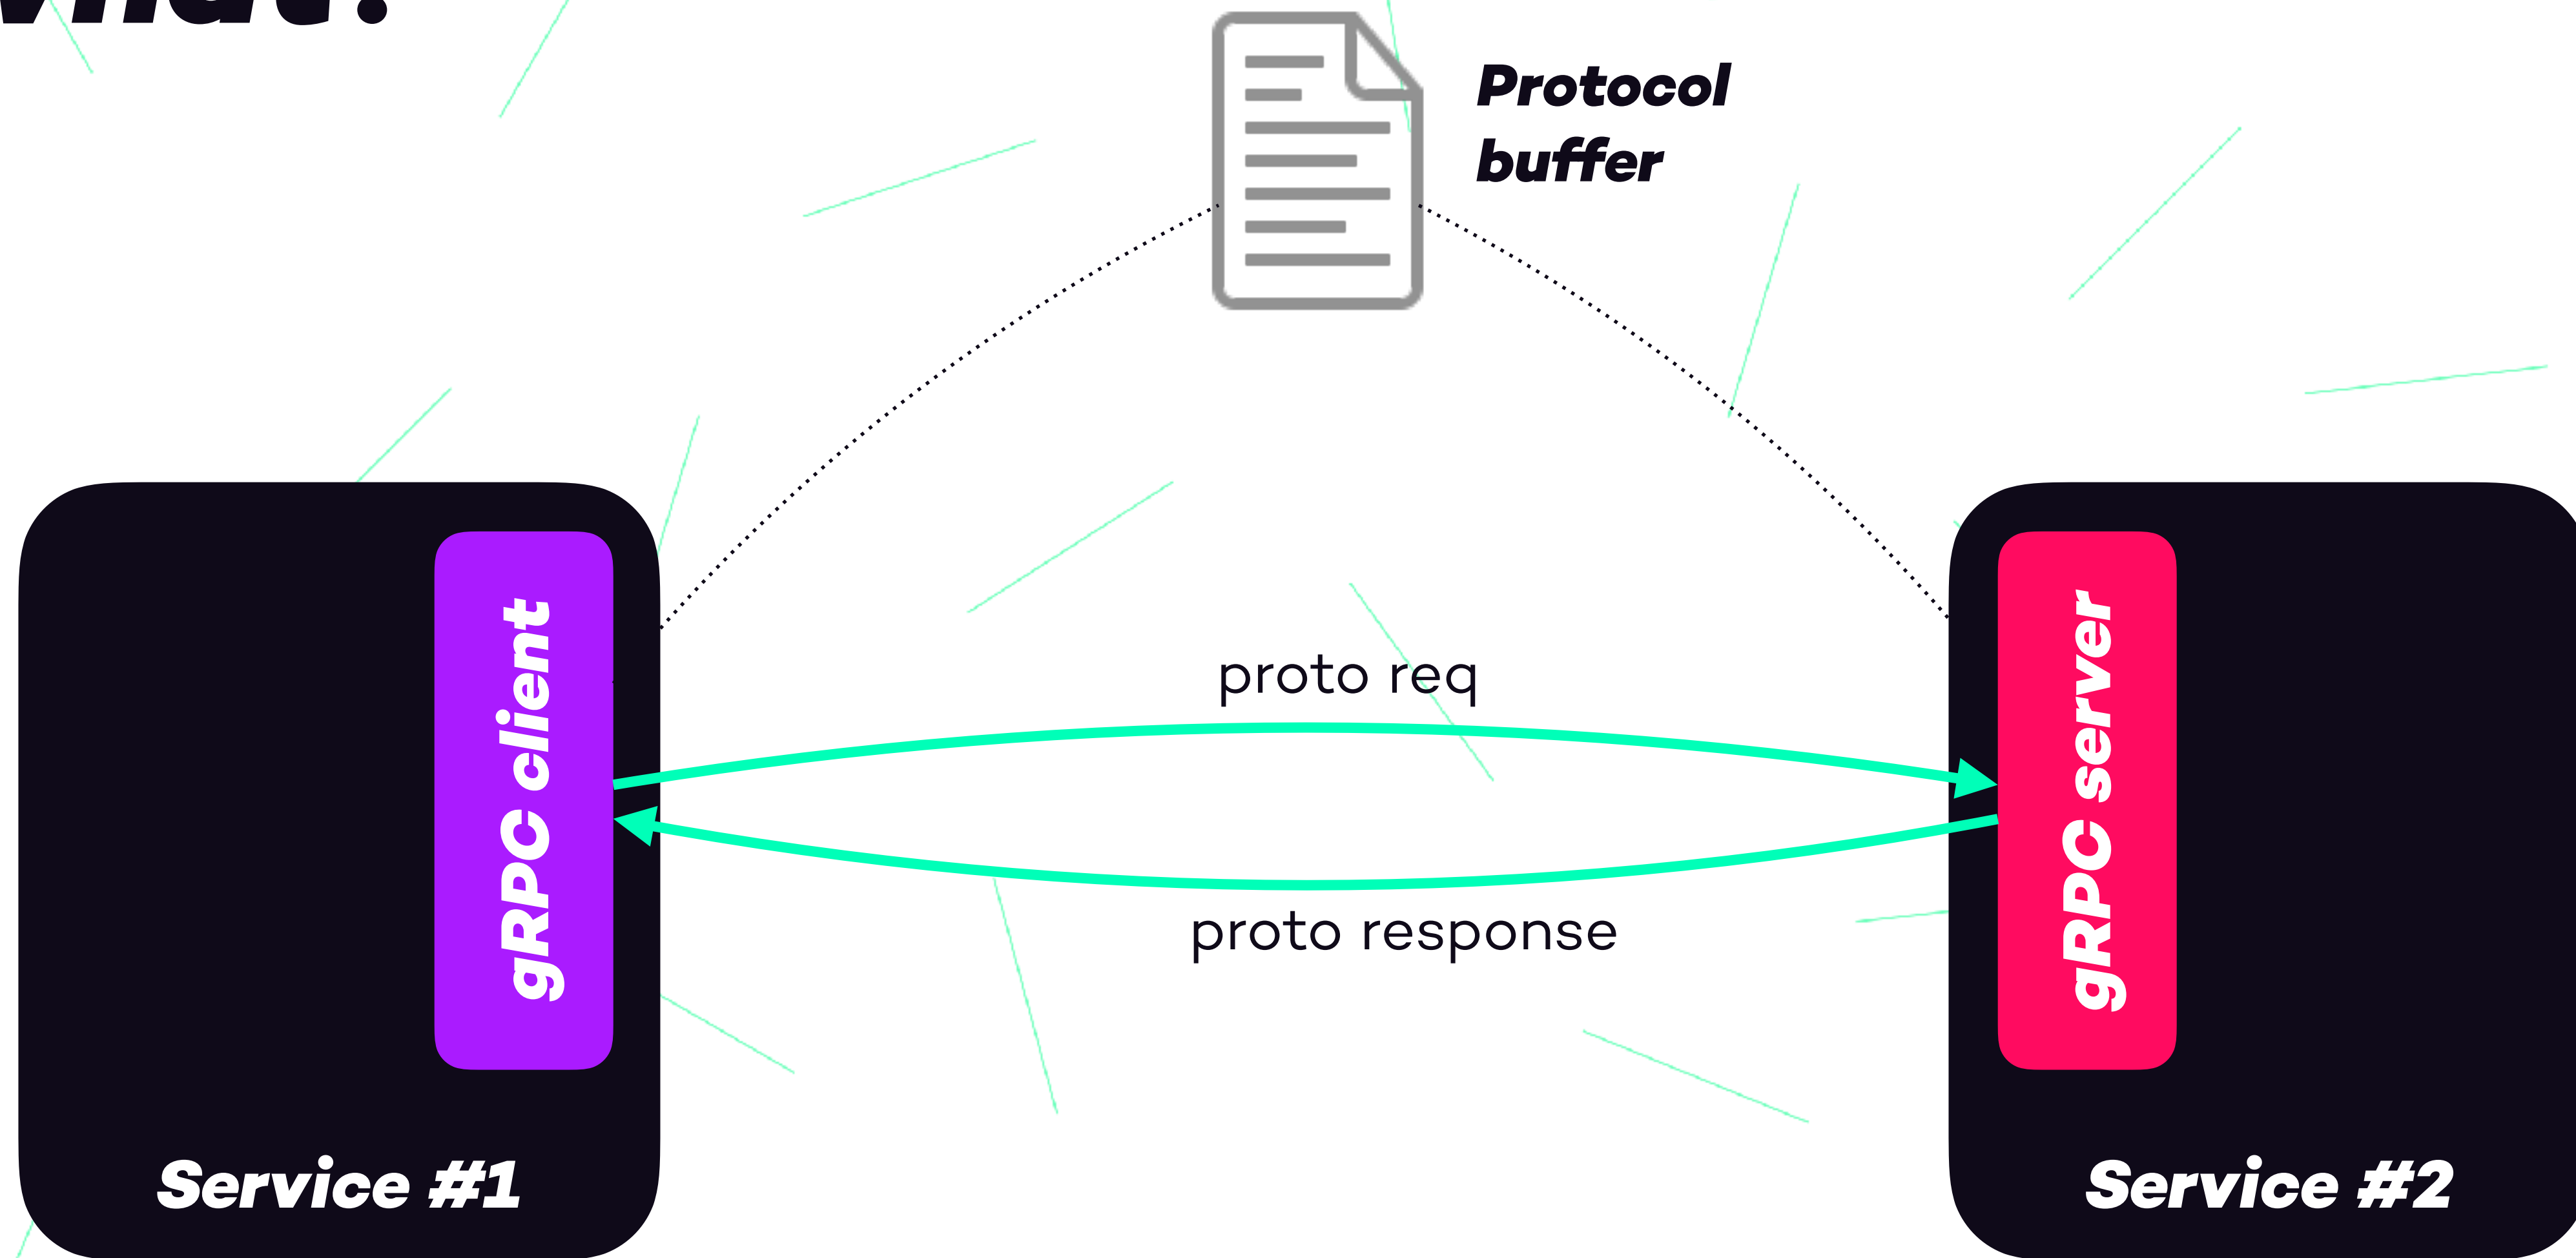
The only thing a service owner has to do in the cluster to make Prometheus scrape their services is to add:

```
annotations:  
  prometheus.io/scrape: 'true'
```

*Service to service communication with*  
**gRPC**

**gRPC**

# What?




# ***What do we use it for?***

## ***It will be our default choice for synchronous calls***

Synchronous service to service communication will be aligned on gRPC

## ***Internal support system will use it to fetch data from our services***

Our internal support system needs information from different services on demand, the service will use gRPC to fetch the data



# ***Why gRPC?***

## ***Simple service definition using Protocol buffers***

Simpler request handling, no need for serialization and deserialization

## ***Binary protocol***

Less overhead in communication.

## ***Works across multiple languages and platforms***

gRPC has a widespread support for multiple languages, making it a perfect fit in our current polyglot architecture

## ***Works great with go and the rest of the ecosystem***

Docker, Kubernetes uses gRPC as long with Go. It's a natural extension for service to service communication. Based on many years of Google experience!



*Back to the*  
***UNICORN...***



# ***Are we there yet?***

***YES! and no..***

***We deploy multiple times a day***

Deployment is autonomous, squads can deploy to production as they please.

***We can easily scale to larger demands, if necessary***

Scaling our infrastructure is easy

***We can tolerate AZ failures to some extend***

Our services are spread across availability zones

***HOWEVER, doing microservices are complex!***

We still need to implement better tracing, using the CNCF project ***OpenTracing*** and Zipkin

We need more insights and smarter routing in our service to service communication, we will be using ***linkerd***.

***Last thing***  
***DO YOU HAVE ANY QUESTIONS?***

# ***Thank you for listening!***

## ***That was it for me!***

If you wanna know more, send me a message in the Cloud Native DK Slack Community.  
Remember to sign up at: <https://cloudnative-dk.herokuapp.com/>

Catch me on Twitter **@phennex**

## ***I will be speaking again:***

- CoDe:U - Continuous Delivery Users Århus (June 20th in INCUBA, Åbogade 15, Aarhus N)
  - Link: <https://www.meetup.com/CoDe-U-AROS/events/239847862/>
- GOTOCon Copenhagen - October 1st
  - Link: <https://gotocph.com/2017/sessions/237>